

# Autonomous Racing Using Reinforcement Learning in ROS2

Adem Jabri

Department of Computer Science, Chair 12  
TU Dortmund University, Germany

22. Januar 2024

# Outline

1 Introduction

2 Objective

3 Methodologie

4 Challenges & Solutions

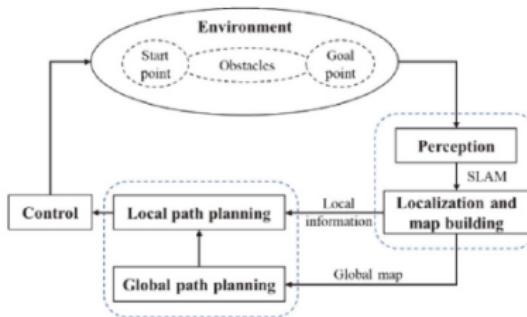
# Motivation



Figure: Autonomous robot systems.

- The key challenge in this field is efficient path following, especially in dynamic and unpredictable environments.
- This thesis explores the use of Reinforcement Learning in ROS2 to address these challenges, offering a novel approach to autonomous navigation.

# Traditional approach



**Figure:** Traditional robot navigation framework.

- **Sensitivity to Conditions:** SLAM can be prone to failure due to changes in the environment or the sensor's conditions.
- **Time-Intensive Map Handling:** The process of establishing and updating the obstacle map in SLAM is time-consuming.
- **Dependence on Sensor Density:** The performance of SLAM algorithms heavily relies on the accuracy of the laser sensor data.

# Traditional approach

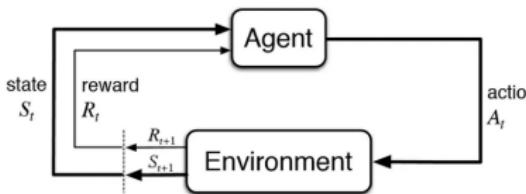


**Figure:** Traditional robot navigation framework.

- **Sensitivity to Conditions:** SLAM can be prone to failure due to changes in the environment or the sensor's conditions.
- **Time-Intensive Map Handling:** The process of establishing and updating the obstacle map in SLAM is time-consuming.
- **Dependence on Sensor Density:** The performance of SLAM algorithms heavily relies on the accuracy of the laser sensor data.

# How will it work with Reinforcement learning?

What is Reinforcement Learning?

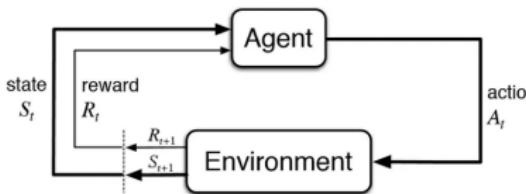


**Figure:** The action-reward feedback loop of a RL model.

- The agent learns to execute tasks by taking actions and observing the outcomes.
- Desirable actions are reinforced through rewards, making them more likely to be repeated.
- Actions that lead to unfavorable outcomes are given negative rewards, or "punishments".
- The goal of reinforcement learning is to find a suitable action model that would maximize the total cumulative reward of the agent.

# How will it work with Reinforcement learning?

What is Reinforcement Learning?



**Figure:** The action-reward feedback loop of a RL model.

- The agent learns to execute tasks by taking actions and observing the outcomes.
- Desirable actions are reinforced through rewards, making them more likely to be repeated.
- Actions that lead to unfavorable outcomes are given negative rewards, or "punishments".
- The goal of reinforcement learning is to find a suitable action model that would maximize the total cumulative reward of the agent.

# How will it work with Reinforcement learning?

- **Interaction-Based Learning:** The RL method enables the robot to learn the optimal navigation policy through direct interaction with the environment.
- **Mapless Navigation:** Unlike SLAM, RL doesn't require pre-existing maps, allowing the robot to navigate without detailed environmental data.
- **Adaptive Policy Development:** Through its learning process, RL develops a navigation policy that guides the robot to its target position adaptively.
- **Strong Learning Ability:** RL can efficiently process complex data and learn from diverse scenarios, improving over time.
- **Sensor Independence:** RL reduces reliance on high sensor accuracy, which is beneficial in environments where sensor data may be unreliable or noisy.

# Outline

1 Introduction

2 Objective

3 Methodologie

4 Challenges & Solutions

# Objective

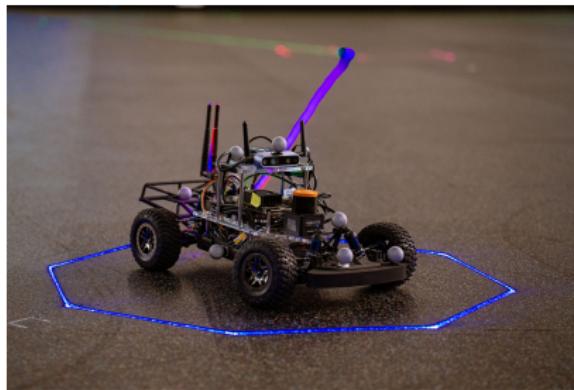


Figure: The autonomous robot.

- Develop an autonomous robot navigation system that can **safely** and **efficiently** navigate **any environment** using a reinforcement learning algorithm and the ROS2 Navigation2 stack.

# Outline

**1** Introduction

**2** Objective

**3** Methodologie

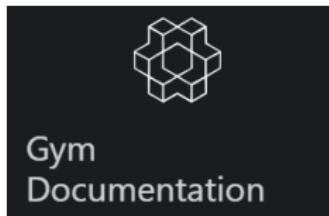
**4** Challenges & Solutions

## Outline

### 3 Methodologie

- Steps
- A Little Bit Into Details

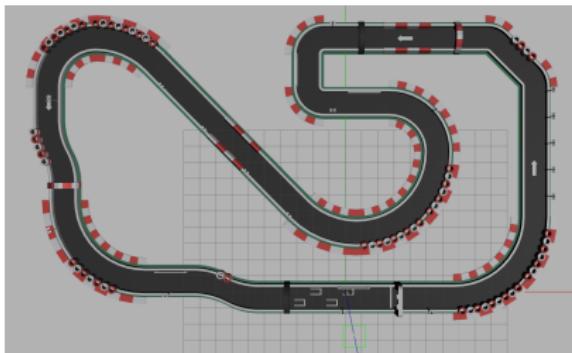
## Step 1:



**Figure:** Integrated tools.

- Implement a Reinforcement Learning algorithm to find the optimal policy for guiding the robot to its target position through interaction with the environment.

## Step 2:



**Figure:** Racetrack in Gazebo Simulator.

- Create a simulation environment using ROS2 and Gazebo, with a robot that is equipped with a LIDAR sensor and a reinforcement learning algorithm.

## Step 3:



Figure: Time to work hard.

- Train the reinforcement learning algorithm to navigate the environment.

## Step 4:



**Figure:** Time to see what you learned.

- Test the robot's ability to navigate the environment safely and efficiently.

## Step 5:



Figure: Truth time.

- Compare the system with the classical navigation algorithms for ROS 2, such as the Navigation2 stack.

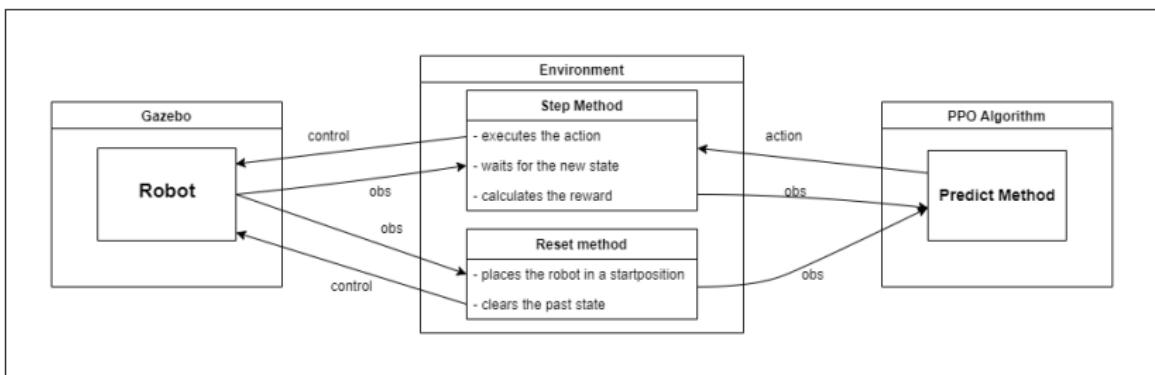
## Outline

### 3 Methodologie

■ Steps

■ A Little Bit Into Details

# RL based navigation system



# Proximal Policy Optimization (PPO)

- A policy describes the behavior of an agent given a state.
- PPO estimates how much better an action is compared to the average action in a given state.
- PPO can be applied to a wide range of environments, both discrete and continuous action spaces.
- It uses a **clipped surrogate objective function**, which prevents too large policy updates, enhancing training stability.
- This objective function provides a balance between **exploration** (trying new things) and **exploitation** (using what is known).

# Proximal Policy Optimization (PPO)

- A policy describes the behavior of an agent given a state.
- PPO estimates how much better an action is compared to the average action in a given state.
- PPO can be applied to a wide range of environments, both discrete and continuous action spaces.
- It uses a **clipped surrogate objective function**, which prevents too large policy updates, enhancing training stability.
- This objective function provides a balance between **exploration** (trying new things) and **exploitation** (using what is known).

# Outline

1 Introduction

2 Objective

3 Methodologie

4 Challenges & Solutions

# Efficient Training



**Figure:** A Race against Time.

- Reinforcement Learning algorithms demand significant computational time.
- Training in a simulated environment makes it faster.
- Implement a reward function that penalizes time to promote quicker navigation.
- Simplify sensor data processing by focusing on the nearest obstacle per sensor section for efficiency.

# Efficient Training



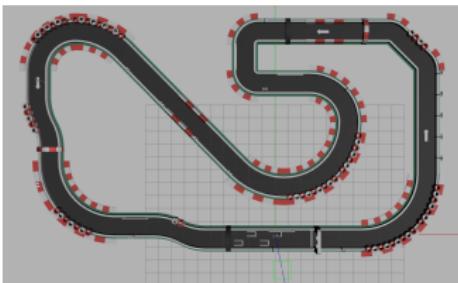
**Figure:** A Race against Time.

- Reinforcement Learning algorithms demand significant computational time.
- Training in a simulated environment makes it faster.
- Implement a reward function that penalizes time to promote quicker navigation.
- Simplify sensor data processing by focusing on the nearest obstacle per sensor section for efficiency.

## Safe Training

- Training in a simulated environment avoids any real-world accidents.
- Including penalties for unsafe actions in the reward function discourages such behaviors.
- Define a safety threshold distance from obstacles, beyond which the robot is considered to have crashed and incorporate it into the reward function.

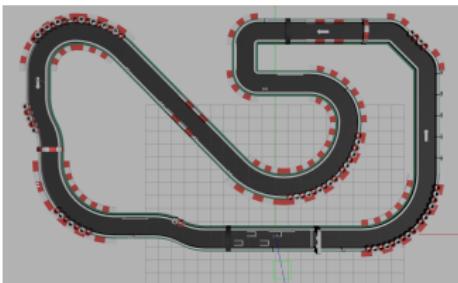
# Overcoming Specialization



**Figure:** Race Track in Gazebo Simulator.

- The model, if overfitted, may navigate effectively only in the specific environment it was trained on, lacking the ability to generalize to new unseen environments.
- Implementing domain randomization to promote generalization.
- Increase the complexity of the environment to progressively challenge the model.

# Overcoming Specialization



**Figure:** Race Track in Gazebo Simulator.

- The model, if overfitted, may navigate effectively only in the specific environment it was trained on, lacking the ability to generalize to new unseen environments.
- Implementing domain randomization to promote generalization.
- Increase the complexity of the environment to progressively challenge the model.

# Overcoming Specialization

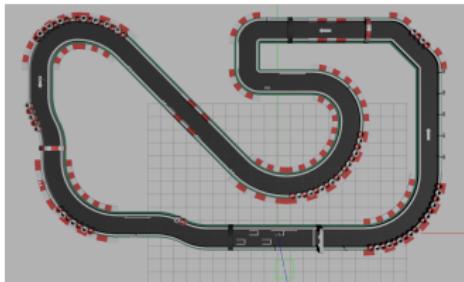


Figure: Original Racetrack.

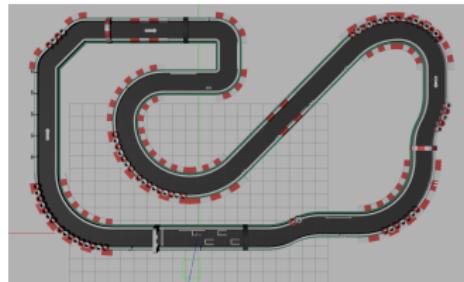


Figure: Mirrored Racetrack.

- Training on both the original and mirrored versions of the racetrack promotes model generalization and mitigates the risk of overfitting, ensuring robust performance in varied track conditions.

## Resources

-  Deep reinforcement learning based mobile robot navigation: A review  
<https://ieeexplore.ieee.org/document/9409758>
-  Proximal Policy Optimization Algorithms  
<https://doi.org/10.48550/arXiv.1707.06347>

# Questions ?